15-122: Principles of Imperative Computation, Summer 1 2012
Assignment 8: The C0VM

William Lovas (`wlovas@cs`)   Tom Cortina (`tcortina@cs`)
Frank Pfenning (`fp@cs`)

Out: Thursday, June 21
Due: Thursday, June 28

1 Overview

In this assignment you will implement a virtual machine for C0, the C0VM. It has been influenced by the Java Virtual Machine (JVM) and the LLVM, a low-level virtual machine for compiler backends. We kept its definition much simpler than the JVM, following the design of C0. Bytecode verification, one of the cornerstones of the JVM design, fell victim to this simplification so in this way the machine bears a closer resemblance to the LLVM. Nevertheless, it is a fully functional design and should be able to execute arbitrary C0 code.

The purpose of this assignment is to give you practice in writing C programs in the kind of application where C is indeed often used in practice. C is appropriate here because a virtual machine has to perform some low-level data and memory manipulation that is difficult to make simultaneously efficient and safe in a higher-level language. It should also help you gain a deeper understanding how C0 (and, by extension, C) programs are executed. We also hope that actual implementations may be useful for further instances of this course, because virtual machines are easy to instrument for debugging purposes, allowing tracing, single-stepping, setting breakpoints, call-counting, etc.

The C0VM is defined in stages, and we have test programs which exercise only part of the specification. We strongly recommend that you construct your implementation following these stages and debug and test each stage before moving on to the next. Each part has its own challenges, but each part should be relatively small and self-contained.

This document describes the structure of the C0VM first, then the instruction set (bytecodes) for the C0VM, and then the file format for a C0 program in bytecode form. After this, you will see the tasks you need to perform, step by step. Read this document very carefully as you prepare to do your work.

1

2   The Structure of the C0VM

Compiled code to be executed by the C0 virtual machine is represented in a byte code format, typically stored in a file ending in extension `.bc0` which we call the *bytecode file*. This file contains numerical and string constants as well as byte code for the functions defined in the C0 source. The precise form of this file is specified in Section 4.

2.1   Types

C0 has so-called *small types* `int`, `bool`, `char`, `string`, `t[]`, and `t*`. Values of these types can be passed to or from functions and held in variables. We think of the small types as constituting two classes: the *primitive types* `int`, `bool`, and `char` and the *reference types* `string`, `t[]` and `t*`. All primitive types are conflated in the C0VM and denoted by `w32`, indicating a 32 bit word. Values of primitive types are denoted by x, i, or n. We also conflate all reference types and write *. Values of reference type are denoted by a for *address*. An address uses either 32 bits or 64 bits. On the Linux machines you use, it will be 64 bits.

In addition C0 has *large types* `struct s` which cannot be passed directly, but must be stored in memory. When the C0VM executes a program, it, too, stores values of large type on *its* heap and refers to them by their address. Calculations of how to access struct fields are performed statically by the compiler, which computes proper offsets for each access. In addition, arrays (referenced by values of type `t[]`), memory cells (referenced by values of type `t*`) and strings (referenced by values of type `string`) are stored on the heap.

2.2   Runtime Data Areas

The C0VM defines several data areas that are used during the execution of a program.

2.2.1   The Program Counter

The program counter `pc` holds the address of the program instruction currently being executed. Unless a nonlocal transfer of control occurs (goto, conditional branch, function call or return) it is incremented by the number of bytes in the current instruction before the next instruction is fetched and interpreted.

2.2.2   The Call Stack

The C0VM has a call stack consisting of *frames*, each one containing local variables, a local operand stack, and a return address. The call stack grows when a function is called and shrinks when a function returns, deallocating the frame during the return.

2.2.3 The Operand Stack

The C0VM is a *stack machine*, similar in design to the JVM. This means arithmetic operations and other instructions pop their operands from an operand stack and push their the result back onto the operand stack. This is in contrast with *register machines* where instructions such as arithmetic operations act on a finite set of registers.

2.2.4 The Heap

At runtime, the C0 heap contains C0 strings, C0 arrays (`t[]`) and C0 cells (`t*`). C0 arrays have to store size information so that dynamic array bounds checks can be performed. It is recommended to use the C heap to implement the C0 heap, that is, allocate strings, arrays, and cells directly using calls to `malloc` and `calloc` or their null-checking variants `xmalloc` and `xcalloc` as defined earlier in this course.

Since C0 is designed as a garbage-collected language, you will not be able to free space allocated on behalf of C0 unless you are willing to implement (or use) a garbage collector. We do not view this as a memory leak of the C0VM implementation. On the other hand, temporary data structures required for the C0VM's own operation should be properly freed.

2.2.5 Constant Pools

Numerical constants requiring more than 8 bits and all string constants occurring in the program will be allocated in constant pools, called the *integer pool* and the *string pool*. They never change during program execution.

2.2.6 Function Pools

Functions, either C0 functions defined in a source file or library functions, are kept in pools called the *function pool* and the *native pool*, respectively. Functions in the function pool are stored with their bytecode instructions, while functions in the native pool store an index into a table of C function pointers that the C0VM implementation can dereference and invoke.

2.3 Frames

A *frame* stores data and partial results during the execution of a function body. It holds a return address, which should be the next instruction in the calling function, an array of local variables denoted by $V[0], \ldots, V[num\_vars - 1]$, and an operand stack for computing expression values. At any point during the execution there is a *current frame* as well as a *calling frame*, where the latter is restored when a function returns to its caller.

Since values of small type may occupy 4 bytes or 8 bytes, the data structures implementing frames, including the array of local variables and the operand stack,

allocate 8 bytes for each data value, whether of primitive type or reference type. From the C perspective, things are simplest if these are of type void*, but to clarify intent, we create a type alias c0_value, defined to be void*, for arbitrary C0 values. This strategy requires us to be able to cast ints to pointers and vice versa without loss of information; the C99 standard leaves this behavior implementation-defined, but the Andrew version of gcc implements such casts. We define a macro VAL(x) to cast an int x to a c0_value, and a converse INT(p) to cast a pointer p to an int. For the latter to work, p must have been created with a VAL(x) cast, otherwise the result is unpredictable except for NULL (which returns 0). So we always have INT(VAL(x)) == x for any integer x.

The main function we have provided (see the file c0vm_main.c) performs some simple tests to verify that casts between void* and int do not lose information. If any of these tests fail, C0VM will abort with an appropriate error message.

2.4   Runtime Errors

In order to fully capture the behavior of C0 programs, you must correctly issue errors for things like dereferencing NULL, indexing into an array outside of its bounds, and dividing by zero. Check the C0 Language Reference for details on what kinds of errors you must handle, and then use the following provided functions to issue appropriate error messages:

```
void c0_memory_error(char *err);     // for memory-related errors
void c0_division_error(char *err);   // for division-related errors
```

For unexpected situations that arise while executing bytecode, situations which could indicate a bug in your VM, you may use the standard C library functions abort or assert to abort your program. See Section 5.3 for more details on this distinction.

3   Instruction Set

We group the instructions by type, in order of increasing complexity from the implementation point of view. We recommend implementing them in order and aborting with an appropriate message when an unimplemented instruction is encountered.

3.1   Stack Manipulation

There are three instructions for direct stack manipulation without regard to types.

```
0x59 dup           S, v -> S, v, v
0x57 pop           S, v -> S
0x5F swap          S, v1, v2 -> S, v2, v1
```

3.2 Arithmetic Instructions

Arithmetic operations in C0 are defined using modular arithmetic based on a two's complement signed representation. This does not match your implementation language (C) very well, where the result of *signed* arithmetic overflow is undefined. On the other hand, *unsigned* arithmetic overflow is defined to be modular arithmetic. We therefore recommend casting `int` as `unsigned int`, perform unsigned arithmetic, then casting back. This does not relieve all manual burden to guarantee C0 compliance, but it goes a long way. We recommend a careful reading of the arithmetic operations in the .

For this implementation strategy to be correct, it is important to verify that our C environment does indeed use a two's complement representation and that the C type of `int` has 32 bits. The provided `main` function (see the file `c0vm_main.c`) performs these checks before starting the abstract machine and aborts the execution if necessary.

In the instruction table below (and for subsequent tables), we use `w32` for the type of primitive values and `*` for the type of reference values. Each line has an opcode in hex notation, followed by the operation mnemonic, followed by the effect of the operation, first on the stack, then any other effect.

```
0x60 iadd        S, x:w32, y:w32 -> S, x+y:w32
0x7E iand        S, x:w32, y:w32 -> S, x&y:w32
0x6C idiv        S, x:w32, y:w32 -> S, x/y:w32
0x68 imul        S, x:w32, y:w32 -> S, x*y:w32
0x80 ior         S, x:w32, y:w32 -> S, x|y:w32
0x70 irem        S, x:w32, y:w32 -> S, x%y:w32
0x78 ishl        S, x:w32, y:w32 -> S, x<<y:w32
0x7A ishr        S, x:w32, y:w32 -> S, x>>y:w32
0x64 isub        S, x:w32, y:w32 -> S, x-y:w32
0x82 ixor        S, x:w32, y:w32 -> S, x^y:w32
```

`idiv` and `irem` can raise exceptions (use the provided function `c0_division_error` to generate a message), and `ishl` and `ishr` need to mask their second argument to 5 bits to bridge the gap between the undefined behavior in C and the specified behavior in C0. Please refer to the C0 language specification for important details.

We have omitted negation `-x`, which the compiler can simulate with `0-x`, and bitwise negation `˜x`, which the compiler can simluate with `x^(-1)`.

3.3 Local Variables

We can move data generically between local variables and the stack, because both primitive types and reference types occupy 4 bytes. The instruction operand `<i>` is one byte following the opcode `0x15` or `0x36` in the instruction stream. Because this is the only way to access a local variable, each function can have at most 256 local variables, which includes the function arguments.

```
0x15 vload <i>      S -> S, v                          (v = V[i])
0x36 vstore <i>     S, v -> S                          (V[i] = v)
```

3.4 Constants

We can push constants onto the operand stack. There are three different forms: (1) a constant null which is directly coded into the instruction, (2) a small *signed* (byte-sized) constant <b> which is an instruction operand and must be sign-extended to 32 bits, and (3) a constant stored in the constant pool. For the latter, we distinguish constants of primitive type from those of reference type because they are stored in different pools.

The two constant loading instructions ildc and aldc take two unsigned bytes as operands, which must be combined into an unsigned integer index for the appropriate constant pool. The integer pool stores the constants directly, and the index given to ildc is an index into the integer pool. The string pool is one large array of character strings, each terminated by '\0'. The index given to aldc indicates the position of the first character; its address is therefore of type char* in C and understood by C as a string.

```
0x01 aconst_null   S -> S, null:*
0x10 bipush <b>    S -> S, x:w32     (x = (w32)b, signed)
0x13 ildc <c1,c2>  S -> S, x:w32     (x = int_pool[(c1<<8)|c2])
0x14 aldc <c1,c2>  S -> S, a:*       (a = &string_pool[(c1<<8)|c2])
```

3.5 Control Flow

Each instruction implicitly increments the program counter by the number of bytes making up the instruction. Control flow instructions change this by jumping to another instruction under certain conditions. The addressing is relative to the address of the branch instruction. The offset is a *signed* 16 bit integer that is given as a two-byte operand to the instruction. It must be signed so we can branch backwards in the program. Note that if_cmpeq and if_cmpne can be used to compare either integers or pointers for equality or inequality, whereas the other comparisons only make sense on integers. The nop "no-op" instruction has no effect.

```
0x00 nop                S -> S

0x9F if_cmpeq <o1,o2>   S, v1, v2 -> S       (pc = pc+(o1<<8|o2) if v1 == v2)
0xA0 if_cmpne <o1,o2>   S, v1, v2 -> S       (pc = pc+(o1<<8|o2) if v1 != v2)
0xA1 if_icmplt <o1,o2>  S, x:w32, y:w32 -> S (pc = pc+(o1<<8|o2) if x < y)
0xA2 if_icmpge <o1,o2>  S, x:w32, y:w32 -> S (pc = pc+(o1<<8|o2) if x >= y)
0xA3 if_icmpgt <o1,o2>  S, x:w32, y:w32 -> S (pc = pc+(o1<<8|o2) if x > y)
0xA4 if_icmple <o1,o2>  S, x:w32, y:w32 -> S (pc = pc+(o1<<8|o2) if x <= y)
```

```
0xA7 goto <o1,o2>        S -> S                    (pc = pc + (o1<<8|o2))
```

3.6 Function Calls and Returns

Function calls come in two forms: invoking a C0 function defined in the same bytecode file and invoking a library function defined in C. In either case, generic arguments v1 through vn are passed on the operand stack and consumed. The C0VM implementation must guarantee that the result v is pushed onto the stack when the function returns. For functions returning void, a dummy value is pushed onto the operand stack to provide a uniform interface to functions.

Function information is stored in the function pool or native pool, both of which are addressed by the instruction operand consisting of two bytes, which must be reconstituted into an *unsigned* 16 bit quantity indexing into the appropriate pool.

```
0xB8 invokestatic <c1,c2> S, v1, v2, ..., vn -> S, v
                          (function_pool[c1<<8|c2] = g, g(v1,...,vn) = v)
0xB0 return    ., v -> .   (return v to caller)
```

When invoking a C0 function (instruction invokestatic) we have to preserve the program counter pc as a return address, the current local variable array V and the current operand stack S. This is the information in a *frame* which is pushed onto a global call stack. Then we set the pc to the beginning of the code for the called function g, allocate a new array of local variables, and initialize it with the function arguments from the old operand stack. We also create a new empty operand stack for use in the called function.

When processing a return instruction we restore the pc, the local variable array V and the operand stack S from the last frame on the call stack. We also need to arrange that the return value is popped from the current operand stack and pushed onto the operand stack of the frame we return to. Some temporary data structures may need to be deallocated at this point.

The main function always has index 0 in the function pool and takes 0 arguments. After reading the file, setting up appropriate data structures, etc., your C0VM implementation should start executing byte code at the beginning of this function and at the end print the final value.

In this assignment, you should use two instances of the abstract data type of stacks, one holding frames (the call stack) the other holding operands (the operand stack).[1] Each of them is used just using pushes and pops, as the stack interface

---

[1]There are a number of other plausible implementation paths for the two stacks. One possibility is inspired by the system call stack for languages such as C. In this strategy there is one global array storing return addresses, local variables, and (in the case of the C0VM) also an area to use as the operand stack. This implementation would not use our abstract datatype of stacks, but implement the stack as an array with a so-called stack pointer which is the current top of the stack. Of course, it is not really treated exactly like an abstract stack, because the access to local variables violates the pure stack discipline. Another possibility is to have a single operand stack, shared throughout the execution,

dictates. Remember, your stacks or arrays will have to share data of different types, namely primitive types and reference types.

A reasonable mapping to C implementation types would be `int` for primitive types and `void*` for reference types, implementing respectively the types `w32` and `*` in the description of the instructions. This allows you to store all C0VM values of small type as type `c0_value` (i.e., `void *`) and cast them to the appropriate type for the operations you need to perform on them and to cast the result back when finished.

3.7   Native Function Calls

Native function calls have the same form as C0 function calls, but the two-byte instruction argument indexes into the native pool, rather than the function pool.

```
0xB7 invokenative <c1,c2>  S, v1, v2, ..., vn -> S, v
```

The value `native_pool[c1<<8|c2]` is an index i into a separate runtime structure, the `native_function_table`. From there you retrieve the address of a function g which has type

```
c0_value (*g)(c0_value*);
```

The type `c0_value*` should be read as an array of pointers of type `c0_value`, which indicates generic data. It also returns a value of generic type. In order to call this function you have to construct an array of length n and store arguments `v1` through `vn` at indices 0 through $n-1$, and then invoke the function g on this array. The result has to be pushed back onto the operand stack.

Native function calls do not therefore involve explicitly managed stack frames. Of course, your abstract machine implementation is using the system stack, so when you call the library function, the library function also uses the system stack, rather than any stack managed explicitly by your virtual machine.

The mapping between native library functions and their indices into the native function table is given as a series of `NATIVE_*` macros in the file `c0vm_natives.h`.

3.8   Memory Allocation, Load, and Store

Besides function calls, the trickiest aspect of the C0VM implementation is the management of the C0 runtime heap. There are two basic options. One is to allocate on the C runtime heap (that is available to your C0VM implementation as it runs) one very large array and perform C0 allocations inside this array. The type of references would then be `int`, values denoting array indices. Alternatively, you can satisfy each C0 allocation request separately by allocating a sufficient amount of space on the C runtime heap, using C pointers to implement C0VM references. The latter option is

---

rather than local ones for each frame. This is possible because the C0VM specification requires that when a function returns, its operand stack must be empty.

advantageous for the purpose of calling native functions because the C runtime heap is shared between the running C0VM implementation and the C0 bytecode program that it executes. In either case, the implementation of allocation must take care to initialize all memory requested by C0 to all zeros, as required by the semantics.

```
0xBB new <s>        S -> S, a:*     (*a is now allocated, size <s>)
```

The new <s> instruction allocates memory for holding data of size s, and returns the address of that memory. Here s is an unsigned byte, expressing the memory size in bytes. The data size is computed statically by the C0 compiler. For example, the C0 expression alloc(int) would translate to new 4, while alloc(struct b) would translate to new n, where n is the size of a struct b in memory in bytes, which is always known at compile time.

```
0xBC newarray <s>  S, n:w32 -> S, a:*     (a[0..n) now allocated)
0xBE arraylength   S, a:* -> S, n:w32     (n = \length(a))
```

The newarray <s> instruction allocates memory for an array, each of whose elements has size s. The number n of elements to allocate is passed on the stack since it cannot in general be known at compile time. Unlike C, in C0 array bounds must be checked on accesses to the array, so an array must store its length. It can be retrieved with the arraylength instruction, which is passed an argument on the operand stack which must be a reference to an array.

The layout for arrays must be specified precisely so that native library functions can reliably convert between C0 and the native format. The first 4 bytes contain the number of elements, the next 4 bytes contain the size of each element, followed by the "raw" array, whose size is $n * s$ bytes. See the text file C_IDIOMS.txt for the C idiom to achieve this kind of layout.

Accessing memory is decomposed into address arithmetic and loading from or storing to a computed address. Address arithmetic comes in two forms. We can add a field offset to access a field of a struct, written below as $a + f$ where f, the instruction operand, is an unsigned one-byte quantity. The offset is computed by the C0 compiler. It does not need to be checked, but a must not be null.

```
0x62 aaddf <f>      S, a:* -> S, (a+f):*  (a != NULL; f field offset)
0x63 aadds          S, a:*, i:w32 -> S, (a+s*i):*
                                     (a != NULL, 0 <= i < \length(a))
```

The second form aadds computes the address of an array element. The operand a on the stack must be the address of an array, and the operand i must be a valid index for this array. The C0VM must issue an error message and abort if this is not a valid index, which can be determined from the stored array length; use the provided function c0vm_memory_error to issue this error. We then use the size s stored with the array to compute the address of the ith element. In your implementation you

will have to be careful to account for additional 8 bytes stored for each array in order to obtain the correct address. Note that one `aadds` instruction is necessary for every array access, even if we access the element at index 0.

Once we have calculated an address that holds a value of small type (either a primitive type or a reference type), we can load it from memory or store something at the given location. For primitive types, we use the `imload` instruction and `mstore` instructions, for reference types the `amload` and `amstore` instructions.

```
0x2E imload    S, a:* -> S, x:w32   (x = *a, a != NULL, load 4 bytes)
0x2F amload    S, a:* -> S, b:*     (b = *a, a != NULL, load address)
0x4E imstore   S, a:*, x:w32 -> S   (*a = x, a != NULL, store 4 bytes)
0x4F amstore   S, a:*, b:* -> S     (*a = b, a != NULL, store address)
```

For character arrays, we also need to be able to load an individual byte and cast it as a value of C0VM type `w32`, zero-extending it. Zero-extension works here because the original C0 type of such a value must have been `char` whose range is limited to 7 bits. This is done by the `cmload` instruction; `cmstore` performs the opposite, masking the given value of type `w32` to 7 bits.

```
0x34 cmload   S, a:* -> S, x:w32   (x = (w32)(*a), a != NULL, load 1 byte)
0x55 cmstore  S, a:*, x:w32 -> S   (*a = x & 0x7f, a != NULL, store 1 byte)
```

4 Bytecode File Format

The bytecode file, usually with extension `.bc0`, is produced by the `cc0` compiler when invoked with the `-b` or `--bytecode` flag. In order to allow you to easily read bytecode, and also write your own bytecode, the binary file is coded in hexadecimal form, where two-digit bytes are separated by whitespace. In addition, the file may contain comments starting with '#' and extending to the end of the line.

We describe the format as pseudo-structs, where we use the types described below. For multi-byte types, each byte is given separately by two hexadecimal digits, with the most significant byte first.

```
u4 - 4 byte unsigned integer
u2 - 2 byte unsigned integer
u1 - 1 byte unsigned integer
i4 - 4 byte signed (two's complement) integer
fi - struct function_info, defined below
ni - struct native_info, defined below
```

The size of some arrays is variable, depending on earlier fields. These are only arrays conceptually, of course. In the file, all the information is just stored as sequences of bytes separated by whitespace.

```
struct bc0_file {
  u4 magic;                          # magic number, always 0xc0c0ffee
  u2 version;                        # version number, currently 2
  u2 int_count;                      # number of integer constants
  i4 int_pool[int_count];            # integer constants
  u2 string_count;                   # number of characters in string pool
  u1 string_pool[string_count];      # adjacent '\0'-terminated strings
  u2 function_count;                 # number of functions
  fi function_pool[function_count];  # function info
  u2 native_count;                   # number of native (library) functions
  ni native_pool[native_count];      # native function info
};

struct function_info {
  u2 num_args;                       # number of arguments, V[0..num_args)
  u2 num_vars;                       # number of variables, V[0..num_vars)
  u2 code_length;                    # number of bytes of bytecode
  u1 code[code_length];              # bytecode
};

struct native_info {
  u2 num_args;                       # number of arguments, V[0..num_args)
  u2 function_table_index;           # index into table of library functions
};
```

We are providing code that reads bytecode files and marshals the information into similar internal C structures.

5 Programming Tasks and Coding Advice

There are many complexities in implementing a virtual machine, especially one that is rich enough so it can execute all of C0! Fortunately, some of the complexities (such as parsing the bytecode file) are taken care of by code we are providing, but others remain. You will complete the code in c0vm.c.

The following are suggested strategies to help you work effectively throughout this project.

5.1 Testing

We have provided a few test cases in test/, but it is more effective to write your own. Write a small file, say test.c0 and the compile it with cc0 -b test.c0, which will create a bytecode file test.bc0. Then run it with ./c0vm test.bc0. Compare your answers with the ones you get with cc0 -o test test.c0 and ./test.

5.2 Incremental Implementation

Implement a subset of the instruction set and test your C0VM implementation on code that only uses the subset. Generate some test cases using `cc0 -b` from simple C0 sources, or use some of the supplied examples that use limited instructions. You should recognize instructions that are valid but not in your subset and give a "*not yet implemented*" message and returning rather than aborting in the same way as for other errors. Test one stage thoroughly before moving on. After extending the machine, first make sure the old, simple examples still run correctly, a process called *regression testing*. The stages follow our discussion of the instruction set:

Task 1 (10 pts) Initialize the following variables correctly in the `execute` function in the `c0vm.c` file: `callStack`, `main_fn`, `V`, `S`, `P`, `pc`.

> **Handin:** No handin is required for this task. It will be included in all of your subsequent handins.

Task 2 (25 pts) [3.1,3.2] Add code to handle arithmetic instructions, plus `bipush`, `swap`, and `return`. (**Note:** The return instruction need only simulate returning from `main` for now.) C0 programs with only a `main` function returning an expression made of small constants can be used to test these capabilities, e.g.,

```
int main() {
  return 15 * ((1<<10) - 24) + 122;
}
```

> **Handin:** When you complete this task and have tested thoroughly, you should hand in this version of the C0VM under the name `c0vm-arithmetic.c`. You may also hand in any bytecode (`.bc0`) files you used to test your code.

Task 3 (10 pts) [3.3,3.4] Add code to deal with local variables and constants. C0 source files containing straight-line code using variables and large constants can be used to test these capabilities, e.g.,:

```
int main() {
  int x = 15122;
  int y = x * x;
  return y;
}
```

> **Handin:** When you complete this task and have tested thoroughly, you should hand in this version of the C0VM under the name `c0vm-locals.c`. You may also hand in any bytecode (`.bc0`) files you used to test your code.

Task 4 (25 pts) [3.5] Add code to handle `goto` and conditionals (e.g., `if_icmpge`). Now you should be able to execute loops, as in

```
int main () {
  int i; int sum = 0;
  for (i = 15; i <= 122; i++)
    sum += i;
  return sum;
}
```

Handin: When you complete this task and have tested thoroughly, you should hand in this version of the C0VM under the name `c0vm-control.c`. You may also hand in any bytecode (`.bc0`) files you used to test your code.

Task 5 (20 pts) [3.6,3.7] Add function calls (`invokestatic`, `invokenative`). This requires managing a call stack in some form, and you will need to revisit `return`. You may want to focus on ordinary C0 function calls (`invokestatic`, `return`) before moving on to native function calls (`invokenative`).

Now your `main` function can call auxiliary functions, such as the ubiquitous recursive factorial function, and library functions that print output:

```
int factorial(int n) {
  if (n == 0) return 1;
  else return n * factorial(n-1);
}
int main () {
  printint(factorial(15));
  println(" is the factorial of 15");
  return 0;
}
```

Handin: When you complete this task and have tested thoroughly, you should hand in this version of the C0VM under the name `c0vm-functions.c`. You may also hand in any bytecode (`.bc0`) files you used to test your code.

Task 6 (10 pts) [3.8] Add the C0 heap, where arrays and structs are allocated. After this, you should be able to run any C0 code, except code that uses unsupported libraries such as `-l img`.

Handin: When you complete this task and have tested thoroughly, you should hand in this version of the C0VM under the name `c0vm-memory.c`. You may also hand in any bytecode (`.bc0`) files you used to test your code.

5.3 Assertions

Ideally, we would establish invariants of the bytecode that we read from a file to make sure no runtime memory or type error occurs. In the JVM this is referred to

as *bytecode verification*. Unfortunately, the current bytecode format does not provide enough information to do this. Even if it did, it would be a major project in itself. So you have to fall back on dynamic checks. These checks come in two categories:

1. The usual checks on the runtime structure of your own code, verifying that pointers are not null, etc.

2. Checks that the C0 bytecode you have read in behaves properly.

Some of the checks in the second category are mandated:

(a) The C0 program must not dereference the C0 null pointer or perform pointer arithmetic on it.

(b) The C0 program must not access memory outside the bounds of a C0 array.

(c) The C0 program must not perform illegal integer division (division by 0, or the min int divided by -1).

If you encounter these runtime errors, you should produce error messages using the provided functions `void c0_memory_error(char *err)` (for memory-related errors) and `void c0_division_error(char *err)` (for division-related errors). The error messages for these should make it clear that this is a runtime error in the bytecode you are executing and not a bug in your machine.

   The first category of checks should in principle be redundant. For example, the `cc0` compiler should never produce a bytecode file that jumps to an invalid address. Nevertheless, bytecode written by hand or a bug in the `cc0` compiler or your VM could lead to such issues. Since C does not guarantee detection of such incorrect jumps or accesses, your code should do that using appropriate `assert` statements, or `ASSERT`, `REQUIRES`, and `ENSURES`. Then, if the bytecode itself or your virtual machine implementation has a bug, it will be discovered as soon as unexpected incorrect behavior occurs. The macro annotations are recommended so that there no undue overhead for correct code when your machine has been debugged.

5.4   The Proliferation of Types

One common issue in writing virtual machines and similar (almost) self-referential code is the proliferation of different types at different levels, sometimes with the same name. It will be important for you to understand the various level of types involved, whether they are signed or unsigned, and what they refer to. In addition you will have to cast between types for various operations.

The following table might help. The C type is only recommended, not required.

C0 type	C0VM type	C type (recommended)
int	w32	int
bool	w32	int
char	w32	int
t[]	*	void*
t*	*	void*
struct s	(none)	(none)
function	function pool index	struct function_info
library fn	native pool index	struct native_info

Sometimes you will have to cast the C type that represent the C0VM types to a form where they are appropriate for the operation to be performed on them. An example of this are the arithmetic operations, which are reliably performed on unsigned ints so we have to cast to and from these types.

5.5   Manage Your Time Well

Remember that this homework is worth **100 points**. You should plan on working on this for an hour or two every day, so you can ask for help early on if you need it. **Don't wait until the last few days!** Post general questions on the bulletin board (e.g., questions about the C0VM specification, wording of tasks, requirements for handin, etc.).

A C0VM Instruction Reference

```
S = operand stack
V = local variable array, V[0..num_vars)

Instruction operands:
<i> = local variable index (unsigned)
<b> = byte (signed)
<s> = element size in bytes (unsigned)
<f> = field offset in struct in bytes (unsigned)
<c> = <c1,c2> = pool index = (c1<<8|c2) (unsigned)
<o> = <o1,o2> = pc offset = (o1<<8|o2) (signed)

Stack operands:
a : * = address ("reference")
x, i, n : w32 = 32 bit word representing an int, bool, or char ("primitive")
v = arbitrary value (v:* or v:w32)


Stack operations

0x59 dup          S, v -> S, v, v
0x57 pop          S, v -> S
0x5F swap         S, v1, v2 -> S, v2, v1

Arithmetic

0x60 iadd         S, x:w32, y:w32 -> S, x+y:w32
0x7E iand         S, x:w32, y:w32 -> S, x&y:w32
0x6C idiv         S, x:w32, y:w32 -> S, x/y:w32
0x68 imul         S, x:w32, y:w32 -> S, x*y:w32
0x80 ior          S, x:w32, y:w32 -> S, x|y:w32
0x70 irem         S, x:w32, y:w32 -> S, x%y:w32
0x78 ishl         S, x:w32, y:w32 -> S, x<<y:w32
0x7A ishr         S, x:w32, y:w32 -> S, x>>y:w32
0x64 isub         S, x:w32, y:w32 -> S, x-y:w32
0x82 ixor         S, x:w32, y:w32 -> S, x^y:w32

Local Variables

0x15 vload <i>    S -> S, v        v = V[i]
0x36 vstore <i>   S, v -> S        V[i] = v
```

```
Constants

0x01 aconst_null   S -> S, null:*
0x10 bipush <b>    S -> S, x:w32      (x = (w32)b, signed)
0x13 ildc <c1,c2>  S -> S, x:w32      (x = int_pool[(c1<<8)|c2])
0x14 aldc <c1,c2>  S -> S, a:*        (a = &string_pool[(c1<<8)|c2])


Control Flow

0x00 nop                  S -> S
0x9F if_cmpeq <o1,o2>    S, v1, v2 -> S        (pc = pc+(o1<<8|o2) if v1 == v2)
0xA0 if_cmpne <o1,o2>    S, v1, v2 -> S        (pc = pc+(o1<<8|o2) if v1 != v2)
0xA1 if_icmplt <o1,o2>   S, x:w32, y:w32 -> S  (pc = pc+(o1<<8|o2) if x < y)
0xA2 if_icmpge <o1,o2>   S, x:w32, y:w32 -> S  (pc = pc+(o1<<8|o2) if x >= y)
0xA3 if_icmpgt <o1,o2>   S, x:w32, y:w32 -> S  (pc = pc+(o1<<8|o2) if x > y)
0xA4 if_icmple <o1,o2>   S, x:w32, y:w32 -> S  (pc = pc+(o1<<8|o2) if x <= y)
0xA7 goto <o1,o2>        S -> S                (pc = pc + (o1<<8|o2))


Functions

0xB8 invokestatic <c1,c2> S, v1, v2, ..., vn -> S, v
                          (function_pool[c1<<8|c2] => g, g(v1,...,vn) = v)
0xB0 return   ., v -> .   (return v to caller)
0xB7 invokenative <c1,c2> S, v1, v2, ..., vn -> S, v
                          (native_pool[c1<<8|c2] => g, g(v1,...,vn) = v)


Memory

0xBB new <s>       S -> S, a:*          (*a is now allocated, size <s>)
0xBC newarray <s>  S, n:w32 -> S, a:*   (a[0..n) now allocated)
0xBE arraylength   S, a:* -> S, n:w32   (n = \length(a))

0x62 aaddf <f>     S, a:* -> S, (a+f):*   (a != NULL; f field offset)
0x63 aadds         S, a:*, i:w32 -> S, (a+s*i):*
                                        (a != NULL, 0 <= i < \length(a))

0x2E imload    S, a:* -> S, x:w32  (x = *a, a != NULL, load 4 bytes)
0x2F amload    S, a:* -> S, b:*    (b = *a, a != NULL, load address)
0x4E imstore   S, a:*, x:w32 -> S  (*a = x, a != NULL, store 4 bytes)
0x4F amstore   S, a:*, b:* -> S    (*a = b, a != NULL, store address)

0x34 cmload    S, a:* -> S, x:w32  (x = (w32)(*a), a != NULL, load 1 byte)
0x55 cmstore   S, a:*, x:w32 -> S  (*a = x & 0x7f, a != NULL, store 1 byte)
```